

Morphology-matched forward readout gates for Tau Core gravity: extending SPARC residual Papers 1–3 from inverse diagnostics to predeclared family tests

Jozsef Olcsak

2026

Abstract

Paper 1 found that residual-blind external structural disturbance labels are associated with increased low-acceleration residual scatter in SPARC galaxies. Paper 2 reversed the question and showed that fixed residual-shape features can infer those disturbance labels better than shuffled controls, while remaining non-unique with respect to MOND-simple and empirical RAR-like baselines. Paper 3 then built a candidate/control framework and introduced an endpoint-conditioned required- S_τ diagnostic: what preparation/readout factor would be needed to absorb a measured residual? This Paper 8 proposal turns that inverse diagnostic into a forward validation gate. The proposed gate predeclares a residual-blind morphology class K_g , maps it to a Tau Core morphology-conditioned readout shell, computes a forward δg^K without using required S_τ as a predictor, and compares matched-family residual structure against wrong-family, shuffled-label, Newtonian baryonic, MOND-simple, and empirical RAR controls. The permitted claim is methodological: the morphology-matched forward-readout gate is a falsifiable, Tau-Core-specific next endpoint for the Paper 1–3 program. This paper does not claim that Tau Core has been empirically validated, that a universal galaxy law has been derived, or that the matched-family endpoint has already beaten MOND/RAR on real SPARC data.

1 Purpose and claim boundary

The first three SPARC residual papers define a cautious progression. Paper 1 asks whether external structural disturbance is associated with residual scatter [2]. Paper 2 asks whether residual-shape features can recover external disturbance labels better than shuffled labels [3]. Paper 3 turns the association into a candidate/control framework and asks what S_τ would be required to absorb the observed residual [1].

That Paper 3 S_τ diagnostic is useful but partly inverse. It asks:

given the measured endpoint, what $S_\tau(R)$ would absorb it?

The present paper defines the next gate:

residual-blind morphology class K_g
→ forward readout δg^K before endpoint scoring.

The main claim is deliberately narrow:

The Tau Core gravity bridge can sharpen Papers 1–3 by replacing inverse required- S_τ residual absorption with a morphology-matched forward-readout candidate/control gate.

This paper does not claim:

- that Tau Core is proven;
- that the morphology-matched gate has already won on real SPARC endpoints;
- that MOND, RAR, or Newtonian baryonic baselines have been superseded;
- that non-axisymmetric morphology classes can be fully tested from one-dimensional rotation curves alone;
- that the morphology families below are the final physical taxonomy.

2 From Paper 3 inverse diagnostics to forward readout

Paper 3 used an endpoint-conditioned diagnostic of the form

$$S_{\tau,\text{req}}(R) = \frac{V_{\text{obs}}(R)/V_N(R) - 1}{\alpha \ln(1 + a_0/a_N(R))}.$$

This quantity is descriptive. Because it uses V_{obs} , it cannot be treated as an independent predictor. Its correct role is to identify what kind of future nonleaky rule would be required.

The forward replacement is:

$$K_g \longrightarrow \mathcal{F}_{K_g} \longrightarrow \delta g^{K_g}(R) \longrightarrow \text{Score}(g_{\text{Newt}} + \delta g^{K_g}).$$

Here K_g is a morphology label assigned before endpoint scoring, $\mathcal{F}_{K_g}$ is the predeclared readout family, and δg^{K_g} is computed without using $S_{\tau,\text{req}}$, endpoint residual gain, or post-hoc family selection.

3 MORPHOLOGY-MATCHED-FORWARD-READOUT-GATE

For a galaxy g , let

$$K_g \in \mathcal{K}$$

be a residual-blind morphology label. Let $\mathcal{F}_K$ denote the Tau Core readout shell assigned to morphology class K . Let

$$R_{g,K}$$

be the residual score after applying the forward readout family $\mathcal{F}_K$ to galaxy g under a frozen geometry and amplitude policy.

The matched-family score is

$$R_{g,\text{matched}} := R_{g,K_g}.$$

The wrong-family score is

$$R_{g,\text{wrong}} := \frac{1}{|\mathcal{K}| - 1} \sum_{K \neq K_g} R_{g,K}.$$

The central paper-level endpoint can be written as

$$\Delta_{\text{matched}} = \text{RMS}_g(R_{g,\text{wrong}}) - \text{RMS}_g(R_{g,\text{matched}}).$$

An equivalent rank endpoint is

$$\text{rank}_g(K_g) = \text{rank}(R_{g,K_g} \text{ among } \{R_{g,K} : K \in \mathcal{K}\}).$$

The gate passes only if the predeclared matched family improves residual structure more than:

1. wrong morphology families;
2. shuffled morphology labels;
3. Newtonian baryonic baseline;
4. MOND-simple baseline;
5. empirical RAR-like baseline.

This is Tau-Core-specific in a way that ordinary MOND/RAR comparisons are not: the discriminating claim is not merely lower residuals, but morphology-indexed readout specificity.

MORPHOLOGY-MATCHED-FORWARD-READOUT-GATE

- 
1. residual-blind morphology label
 2. formula-shell selection
 3. geometry and amplitude discipline
 4. matched-vs-wrong family endpoint
 5. shuffled-K null
 6. baseline comparison

Figure 1: The proposed morphology-matched forward-readout gate. The labels and family choices must be fixed before endpoint scoring.

4 Morphology-family registry

The current gravity bridge centralizes a morphology-conditioned TGP-like readout lane. The corresponding first-pass family registry is not a universal galaxy law. It is a set of candidate readout shells whose selection must be controlled by residual-blind morphology labels.

Table 1: Claim-bounded morphology-family registry for the forward-readout gate.

Family	Morphology label	Formula shell	First-pass status
K_{tail}	Scale-tail spiral / diffuse LSB disk	$\sigma_{\text{morph}}(R) \sim A_K / (R + R_c)^2$	1D testable
K_{exp}	Regular exponential disk	$\sigma_{\text{morph}}(R) \sim A_K e^{-R/R_d}$	1D testable
K_{compact}	Compact finite support	$I_{\text{morph}}(R) \rightarrow I_\infty$ and $\delta g_R \sim -I_\infty / R^2$	1D testable
K_{ring}	Ring or resonance feature	localized $\sigma_{\text{morph}}(R)$ near R_{ring}	1D proxy
K_{flare}	Thick or flared disk	projected thickness/flaring kernel	1D proxy
$K_{m=2}$	Barred spiral / $m = 2$ mode	harmonic $\sigma_{\text{morph}}(R, \phi)$	velocity-field preferred
$K_{m=1}$	Lopsided / $m = 1$ mode	harmonic $\sigma_{\text{morph}}(R, \phi)$	velocity-field preferred




Figure 2: First-pass observability split for the current family registry. SPARC one-dimensional rotation curves are most suitable for axisymmetric or radially localized families; barred and lopsided branches require velocity-field follow-up for a full test.

5 Candidate/control continuity from Paper 3

The gate preserves the Paper 3 candidate/control discipline. It does not discard difficult controls after seeing endpoints. The same object may carry a different role under the forward gate:

candidate/control role $\longrightarrow$ predeclared morphology row $\longrightarrow$ matched/wrong/shuffled score.

Table 2: Paper 3 roles carried into the Paper 8 forward-gate protocol.

Galaxy	Paper 3 role	Paper 8 role
DDO126	positive anchor	matched-family candidate after residual-blind morphology labeling
DDO50	quiet control	false-positive suppression control
WLM	geometry stress	observability-caveated endpoint row
DDO168	ladder support	candidate/control consistency row
DDO154	countercontrol	wrong-family and rank-sensitivity stress row

The positive paper-level endpoint is not:

Tau Core fits every galaxy better.

The first meaningful endpoint is:

the correct predeclared morphology family ranks better than wrong families.

This is why shuffled- K nulls are mandatory. Without the shuffled-label null, a morphology-indexed family could accidentally behave like a generic flexible residual smoother.

6 Synthetic protocol fixture

This package includes a synthetic score fixture to verify the endpoint logic. It is not real SPARC evidence. It exists only so the public tests can check that the manuscript, tables, and figure pipeline agree about the intended endpoint.




Figure 3: Synthetic protocol fixture for the endpoint schema. A real pass would require the matched-family score to beat wrong-family and shuffled-label controls, then be compared to Newtonian baryonic, MOND-simple, and empirical RAR baselines.

7 Available-data Tau-proxy preflight

The repository also includes a first available-data preflight of the matched-family idea. This is not the final Paper 8 endpoint, because the family functions are still proxy functions built from the locally available Tau-proxy channel rather than the final morphology-specific δg^K readout shells. The preflight nevertheless tests the important next question: whether a metadata-proxy morphology family assigned before scoring performs better than wrong Tau-proxy families on held-out galaxies.

In the current 175-galaxy point-level packet, family amplitudes are fit on the train split only. On the 44-galaxy holdout split, the morphology-matched proxy beats the wrong-family mean in 0.568 of galaxies, beats the TPG/v6 comparator in 0.636 of galaxies, and beats the MOND comparator in 0.614 of galaxies. A deterministic 1000-permutation shuffled-label null gives a holdout null mean of 0.0133 for the mean matched-minus-wrong endpoint, compared with the observed value -0.0326 , with $p \simeq 0.264$ for shuffled labels being at least as good. Thus the available proxy signal is positive but not yet a discovery-level morphology-specific result. These numbers are recorded as a preparation result, not as empirical validation of Tau Core. They motivate replacing the generic Tau proxy with source-native morphology-family readout shells and a stronger shuffled-label null in the final endpoint.

A second preflight replaces the single available Tau-proxy channel with different radial formula-shell proxies assigned by morphology class. This is closer in form to the intended Paper 8 family test, but the current available-data version is weaker: on the same 44-galaxy holdout split it beats

the wrong-shell mean in 0.500 of galaxies, beats TPG/v6 in 0.477 of galaxies, and has a shuffled-null $p \simeq 0.479$ for the mean matched-minus-wrong endpoint. This negative preflight result is useful. It indicates that merely assigning naive radial shell shapes is not sufficient; the final endpoint needs source-native morphology labels, dimensionally audited 4D readout shells, and a nonleaky amplitude policy.

8 Minimal numerical protocol

The minimal public endpoint protocol is:

1. Assign an external or residual-blind morphology label K_g to each galaxy.
2. Select a formula shell $\mathcal{F}_{K_g}$ from the predeclared registry.
3. Fix geometry parameters from morphology, not from residual endpoints.
4. Use either a frozen amplitude rule or a tightly bounded amplitude fit.
5. Compute the matched-family score.
6. Compute wrong-family scores for all other candidate families.
7. Run shuffled- K nulls.
8. Compare against Newtonian baryonic, MOND-simple, and empirical RAR baselines.

For SPARC-only first-pass work, the cleanest initial family set is:

$$\{K_{\text{tail}}, K_{\text{exp}}, K_{\text{compact}}, K_{\text{ring}}, K_{\text{flare}}\}.$$

The $m = 1$ and $m = 2$ branches should be treated as caveated proxy rows in one-dimensional rotation curves and as real tests only when velocity-field products are available.

9 Why this would be a stronger Tau Core test

The new gate is stronger than a residual association because it creates a family-selection burden. A theory with morphology-indexed readout families should not only reduce residuals; it should reduce them more when the correct morphology family is used than when a wrong family is used. MOND-simple and empirical RAR baselines are strong comparators, but they do not naturally provide a predeclared morphology-family rank test of this kind.

The result that would matter is therefore:

$$\Delta_{\text{matched}} > 0$$

with a rank distribution in which K_g is systematically near the top, under residual-blind morphology labels and with shuffled-label controls. If this fails, the bridge loses specificity. If it passes, the Paper 1–3 residual-disturbance program becomes a forward-predictive morphology-readout program.

10 Readiness and remaining blockers

Table 3: Current readiness of the Paper 8 forward endpoint.

Layer	Status
Paper 1 residual-blind disturbance association	available
Paper 2 residual-shape inference and shuffled-label null	available
Paper 3 candidate/control ladder and required- S_τ diagnostic	available
Morphology-family formula registry	defined in this package
Available-data matched Tau-proxy preflight	run; proxy only
Accepted source-native component tables for full sample	pending
Residual-blind morphology labels for full endpoint sample	pending
Real matched-vs-wrong family endpoint	not yet run

11 Conclusion

Paper 3 is the natural launching point for this forward gate because it already contains the candidate/control framework and the required- S_τ diagnostic. The present contribution is to state the next gate in a falsifiable form:

morphology label $K_g \rightarrow$ formula shell $\mathcal{F}_{K_g} \rightarrow \delta g^{K_g} \rightarrow$ matched vs wrong family endpoint.

The claim-safe breakthrough target is not a universal fit claim. It is a morphology-specific prediction claim: the correct Tau Core morphology readout family, chosen before endpoint scoring, should improve residual structure more than wrong morphology families and shuffled morphology labels, while remaining competitive against Newtonian baryonic, MOND-simple, and empirical RAR baselines.

That is the first clean endpoint where the Tau Core bridge can become more than residual commentary.

References

- [1] Jozsef Olcsak. Sparc projection-sensitive residual candidate paper 3 reproducibility package, 2026.
- [2] Jozsef Olcsak. Sparc residual-disturbance paper 1 reproducibility package, 2026.
- [3] Jozsef Olcsak. Sparc residual-disturbance paper 2 reproducibility package, 2026.